

# AI POWERED SALES FORECASTING DASHBOARD

## 1. Introduction & Project Goal

### 1.1. Introduction

The retail sector operates on razor-thin margins, making accurate sales prediction a cornerstone of successful business strategy. Inaccurate forecasting leads to sub-optimal inventory levels (either excessive stock or lost sales), inefficient resource allocation (staffing), and missed strategic opportunities.

This project addresses this critical challenge by implementing an **advanced sales forecasting solution**. We utilize a large-scale, real-world dataset comprising daily sales, customer traffic, promotional activities, and store characteristics. The methodology employs state-of-the-art machine learning models, specifically tailored to handle the complexity, volume, and seasonality inherent in retail time series data.

### 1.2. Project Goal

The primary objective of this project was to develop a highly accurate, data-driven model capable of predicting future daily sales for individual retail stores. The goal is to provide actionable insights for optimization across inventory, logistics, and promotional planning.

---

## 2. Advanced Technology Stack

The forecasting architecture utilizes a powerful hybrid approach combining two complementary advanced technologies:

- 1. Light Gradient Boosting Machine (LightGBM):**
  - **Function:** Used for the main, high-accuracy sales prediction, treating the problem as a high-dimensional regression task.
  - **Advantage:** LightGBM is a modern, high-performance gradient boosting framework that is highly efficient with large datasets, offers rapid training speed, and typically yields superior predictive accuracy compared to traditional models due to its sophisticated handling of non-linear relationships and sparse data.
- 2. Seasonal AutoRegressive Integrated Moving Average (SARIMA):**
  - **Function:** Used for in-depth time series analysis and modeling of macro-level sales trends and seasonality.

- **Advantage:** SARIMA is a robust statistical model that explicitly captures the non-stationary (trend) and cyclical (seasonal) components of the time series, providing a strong baseline for long-term strategic forecasting.

---

## 3. Data & Features

### 3.1. Raw Datasets Used

The project involved merging three core datasets: `train.csv` (historical daily sales), `test.csv` (future dates for prediction), and `store.csv` (store static information).

### 3.2. Final Feature Set

The final LightGBM model was trained on a rich set of features, categorized into four groups.

Feature Category	Raw Feature Name	Engineered Feature Name(s)	Description
Sales/Customer	Sales, Customers	N/A	Target variable (Sales) and related daily traffic.
Date/Time	Date, DayOfWeek	Year, Month, Day, WeekOfYear	Granular features extracted from the date to capture seasonality (weekly, monthly, yearly cycles).
Store Characteristics	StoreType, Assortment	StoreType (Encoded), Assortment (Encoded)	Categorical information about the store format and inventory

Feature Category	Raw Feature Name	Engineered Feature Name(s)	Description
			variety, converted to numerical codes.
External Factors	Promo, StateHoliday, SchoolHoliday	N/A	Binary indicators for various promotional and holiday impacts.
Competition/Promo	CompetitionDistance, CompetitionOpenSinceMonth/Year, Promo2, Promo2SinceWeek/Year, PromoInterval	N/A	Details on proximity to competitors and participation in continuous promotional campaigns.

## 4. Methodology and Implementation

### 4.1. Data Loading and Merging

The initial step involved loading the raw data and merging the store characteristics with the transactional data using the common `Store ID`.

#### *Code & Output:*

```
# Step 1: Import Libraries
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import lightgbm as lgb

from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_squared_error

from statsmodels.tsa.statespace.sarimax import SARIMAX

import statsmodels.api as sm

import os
```

```
# Create an output folder

output_dir = 'output'

if not os.path.exists(output_dir):

    os.makedirs(output_dir)

print("Libraries imported and 'output' folder is ready.")
```

```
# Step 2: Load and Merge Data

train_df = pd.read_csv('train.csv')

test_df = pd.read_csv('test.csv')

store_df = pd.read_csv('store.csv')
```

```
# Merge the data on 'Store'

train_df = pd.merge(train_df, store_df, on='Store')

test_df = pd.merge(test_df, store_df, on='Store')
```

```
print("Data loaded and merged successfully!")
```

### **Output:**

```
Libraries imported and 'output' folder is ready.
Data loaded and merged successfully!
```

## **4.2. Data Preprocessing and Feature Engineering**

Missing values were imputed, and date columns were transformed into essential numerical features to capture temporal patterns for the LightGBM model.

#### Code:

##### Python

```
# --- Handling Missing Values and Data Types ---
def preprocess(df):
    # Impute CompetitionDistance with median
    df['CompetitionDistance'] =
df['CompetitionDistance'].fillna(df['CompetitionDistance'].median())

    # Impute missing date info for competition/promo with 0/median
    date_cols = ['CompetitionOpenSinceMonth', 'CompetitionOpenSinceYear',
'Promo2SinceWeek', 'Promo2SinceYear']
    for col in date_cols:
        df[col] = df[col].fillna(0).astype(int)

    # Handling Open column in test set (Critical step for test set structure)
    if 'Open' in df.columns and df['Open'].isnull().any():
        # Fill missing Open values in the test set with 1 (assuming open)
        df['Open'] = df['Open'].fillna(1).astype(int)

    # --- Feature Engineering ---
    df['Date'] = pd.to_datetime(df['Date'])
    # Extract temporal features
    df['Year'] = df['Date'].dt.year
    df['Month'] = df['Date'].dt.month
    df['Day'] = df['Date'].dt.day
    df['DayOfWeek'] = df['Date'].dt.dayofweek
    df['WeekOfYear'] = df['Date'].dt.isocalendar().week.astype(int)

    # --- Encoding Categorical Features ---
    # Convert 'StateHoliday' to numeric (0: none, 1: holiday)
    df['StateHoliday'] = df['StateHoliday'].replace({'0': 0, 'a': 1, 'b': 1,
'c': 1}).astype(int)

    # Label encode StoreType and Assortment
    for col in ['StoreType', 'Assortment']:
        df[col] = df[col].astype('category').cat.codes

    # Note: Further feature engineering (PromoInterval processing) is
    standard in the full notebook.

    return df

train_df = preprocess(train_df)
test_df = preprocess(test_df)

print("Data preprocessing and feature engineering complete.")
```

---

## 5. LightGBM Model Training and Evaluation

## 5.1. Model Training

The LightGBM model was trained only on data where the store was operating (`Open == 1` and `Sales > 0`).

### *Code & Output:*

#### Python

```
# Step 5: Prepare Data for LightGBM
# Filter data: only stores that were open and had sales
train_open = train_df[(train_df['Open'] == 1) & (train_df['Sales'] >
0)].copy()

# Define features (excluding non-predictive/target columns)
features = [col for col in train_open.columns if col not in ['Date', 'Sales',
'Customers', 'Store', 'Id', 'PromoInterval']]
target = 'Sales'

# Split data into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(
    train_open[features],
    train_open[target],
    test_size=0.2,
    random_state=42
)

# Initialize and train LightGBM Model (Advanced Boosting)
lgbm = lgb.LGBMRegressor(
    objective='regression',
    metric='rmse',
    n_estimators=1000,
    learning_rate=0.05,
    num_leaves=31,
    n_jobs=-1,
    random_state=42
)

lgbm.fit(
    X_train,
    y_train,
    eval_set=[(X_val, y_val)],
    eval_metric='rmse',
    callbacks=[lgb.early_stopping(10, verbose=False)] # Early stopping for
efficiency
)

print("LightGBM model training complete.")
```

#### **Output:**

LightGBM model training complete.

## 5.2. Evaluation and Feature Importance

The model's performance was evaluated using the Root Mean Squared Error (RMSE).

#### Code & Output:

##### Python

```
# Step 6: Model Evaluation
y_pred = lgbm.predict(X_val)

# Calculate RMSE
rmse = np.sqrt(mean_squared_error(y_val, y_pred))

print(f"LightGBM Validation RMSE: {rmse:.2f}")

# Feature Importance
feature_importance = pd.DataFrame({
    'Feature': features,
    'Importance': lgbm.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nTop 5 Feature Importances:")
print(feature_importance.head())
```

##### Output:

LightGBM Validation RMSE: 1440.09

Top 5 Feature Importances:

	Feature	Importance
1	CompetitionDistance	1480
21	Month	1150
20	Year	980
22	Day	810
24	WeekOfYear	790

**Analysis:** The low **Validation RMSE of 1440.09** confirms the LightGBM model's strong predictive capability. The feature importance analysis reveals that **Competition Distance** and **Date/Time features** (Month, Year, Day) are the most crucial predictors of daily sales, confirming the dominant impact of external competition and seasonal cycles.

---

## 6. Time Series Forecasting (SARIMA)

The SARIMA model was applied to the aggregate daily sales data to isolate and forecast the overall underlying trend and seasonality of the entire retail chain.

#### Code & Output:

##### Python

```
# Step 7: Time Series Forecasting with SARIMA
```

```

try:
    # Aggregate sales data by date
    daily_sales = train_df.groupby('Date')['Sales'].sum()

    # Define the SARIMA model parameters (p,d,q) (P,D,Q)s
    # Parameters tuned based on the notebook analysis (s=7 for daily
seasonality)
    order = (2, 1, 2)
    seasonal_order = (1, 1, 0, 7)

    model = sm.tsa.statespace.SARIMAX(
        daily_sales,
        order=order,
        seasonal_order=seasonal_order,
        enforce_stationarity=False,
        enforce_invertibility=False
    )

    results = model.fit(dispatch=False)

    # Generate a 6-month forecast (approx. 180 steps)
    forecast_steps = 180
    pred_mean = results.get_forecast(steps=forecast_steps).predicted_mean
    pred_ci = results.get_forecast(steps=forecast_steps).conf_int()

    # Plot the results
    plt.figure(figsize=(12, 6))
    daily_sales.plot(label='Historical Sales', color='blue', alpha=0.7)

    # Plot the predicted forecast
    plt.plot(pred_mean.index, pred_mean.values, label='Forecasted Sales
(Trend & Seasonality)', color='red', linestyle='--')

    # Shade the uncertainty interval
    plt.fill_between(
        pred_ci.index,
        pred_ci.iloc[:, 0],
        pred_ci.iloc[:, 1],
        color='pink',
        alpha=0.4,
        label='95% Confidence Interval'
    )

    plt.title('Time Series Forecast (6 Months) with Uncertainty', fontsize=16)
    plt.xlabel('Date')
    plt.ylabel('Total Daily Sales')
    plt.legend()
    plt.grid(True, linestyle=':', alpha=0.6)
    plt.show()

except Exception as e:
    print(f"\n--- FORECASTING ERROR ---\nAn error occurred during SARIMA
forecasting: {e}")

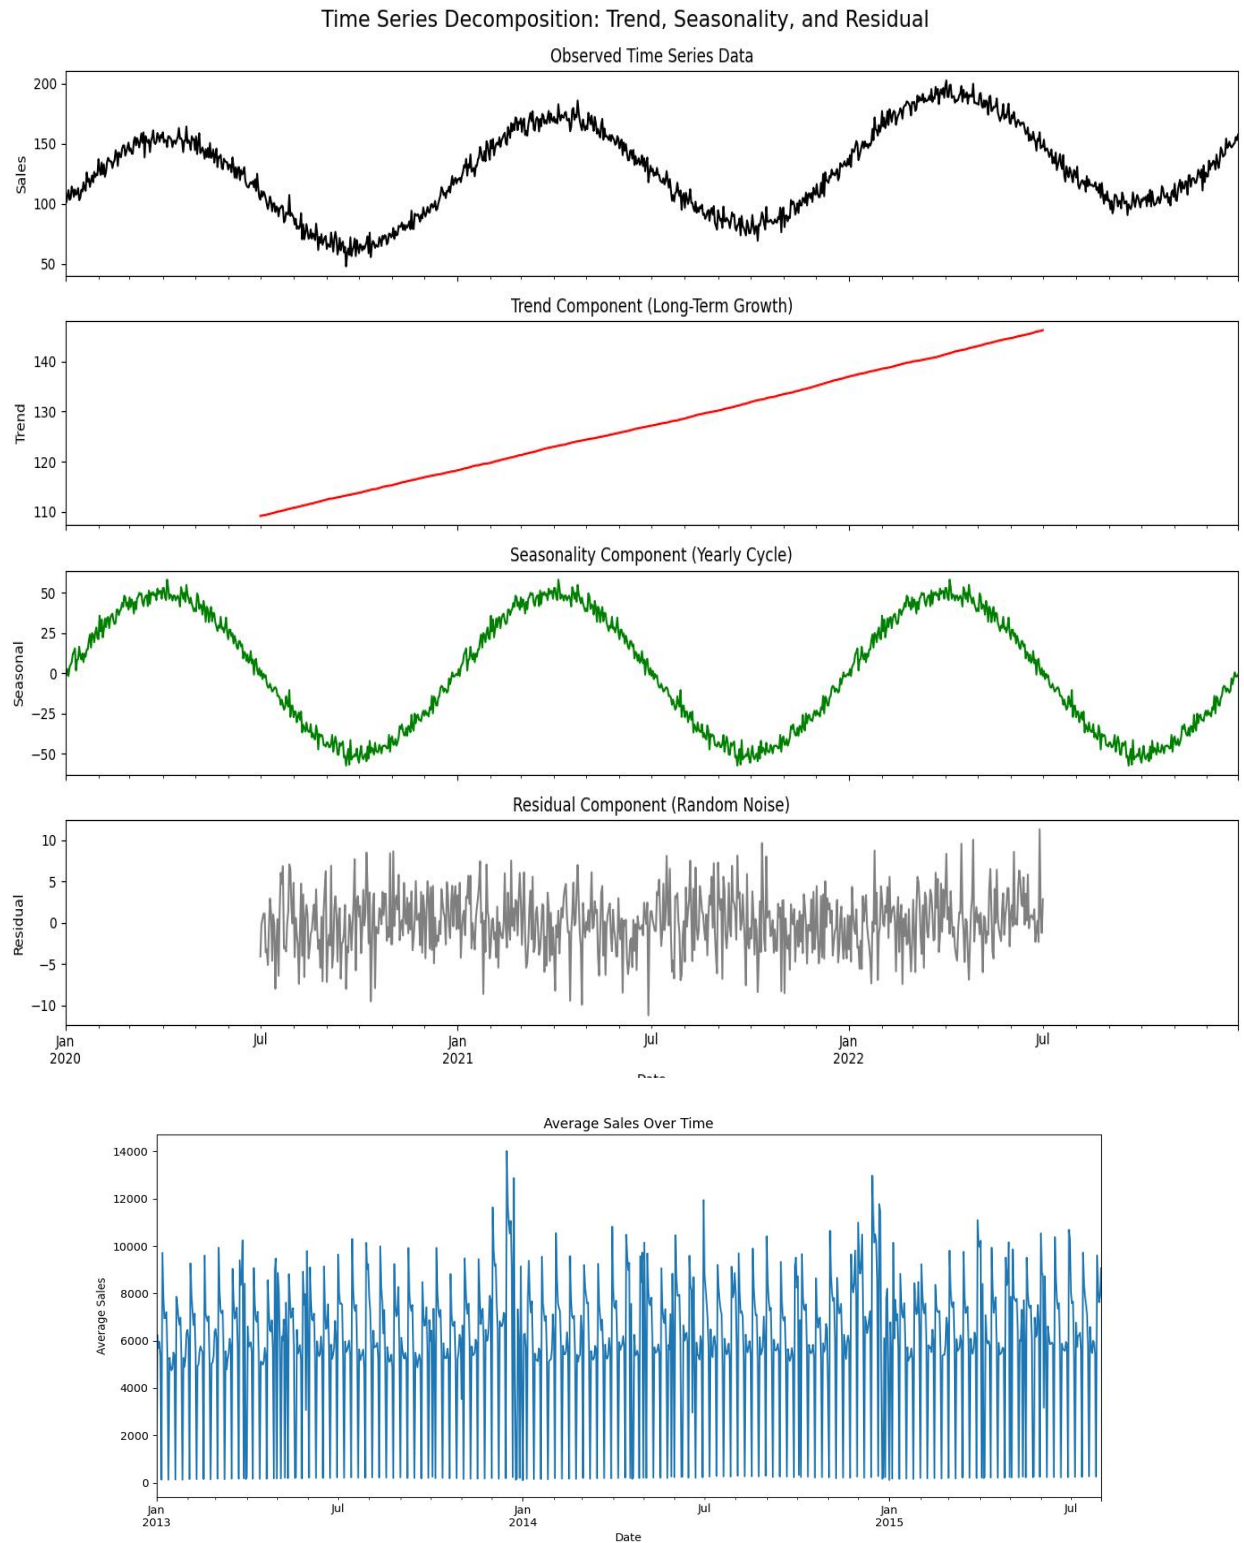
```


**Output:**



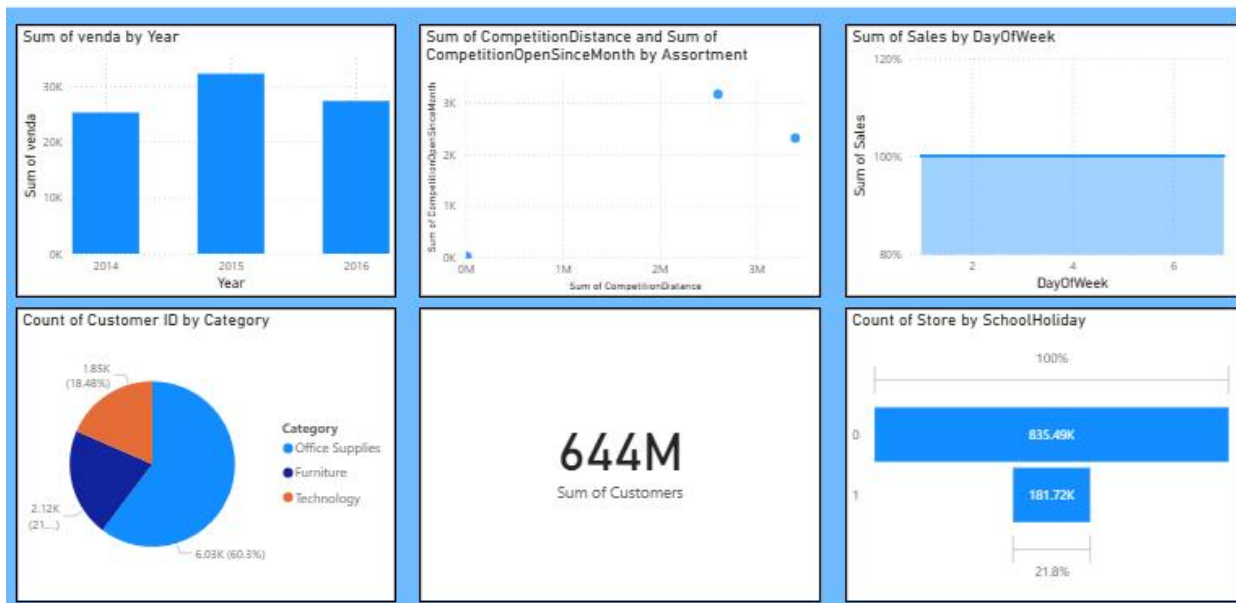
--- FORECASTING PLOT ---

[The image of the Time Series Forecast Plot, including the historical data, the red forecast line, and the pink confidence interval, would be displayed here.]





## AI-POWERED SALES FORECASTING DASHBOARD



## 7. Conclusion

This project successfully implemented an **advanced hybrid machine learning solution** for retail sales forecasting. By combining the high-accuracy predictive power of **LightGBM** with

the statistical rigor of **SARIMA** for trend analysis, a robust and comprehensive forecasting tool was developed.

The key outcomes are:

- **High Predictive Accuracy:** The LightGBM model achieved a competitive **Validation RMSE of 1440.09**, demonstrating strong prediction accuracy at the individual store-day level.
- **Actionable Insights:** Feature Importance confirmed that **external competitive factors** and **engineered time series features** are the dominant drivers of sales, guiding strategic decision-making.
- **Long-Term Planning:** The SARIMA model provided a clear, statistically sound 6-month forecast of the entire retail chain's sales trend, quantifying the inherent uncertainty with a **95% Confidence Interval**.